

ShrinkNets

Guillaume Leclerc, Raul Castro Fernandez, Samuel Madden

Massachusetts Institute of Technology
[leclerc,raulcf,madden]@mit.edu

1 INTRODUCTION

When designing neural networks, one of the key parameters is the network size, i.e., the number of layers and neurons per layer. Choosing these parameters appropriately can dramatically affect performance, yet there is no reliable way to efficiently set them. Although many search strategies and heuristics [?] have been proposed, including random search, meta-gradient descent [?], Gaussian processes [?], and Parzen Estimators [?], these methods generally require a compute-intensive search of parameter space.

In this paper we present a method to automatically find an appropriate network size, drastically reducing optimization time. The key idea is to *learn* the right network size at the same time that the network is learning the main task. For example, for an image classification task, with our approach we can provide the training data to a network—without sizing it a priori—and expect to end up with a network that has learned to classify images with an accuracy similar to a the best manually engineered network. Our approach has two main benefits. First, no longer need to choose a network size before training. Second, the final network size will tuned to be appropriate for the task at hand, and not larger. This is important because over-sized networks have a lower inference throughput and higher memory footprint.

Our approach has two main challenges. First, we need a way to dynamically size the network during training. Second, we need a loss function that optimizes for the additional task of sizing, without deteriorating the learning performance of the main task. Our approach, called ShrinkNets, copes with both challenges.

2 SHRINKNETS

During training, our approach starts with an explicitly over-sized network. As training progresses, we learn which neurons are not contributing to learning and remove them dynamically, effectively shrinking the network. This method requires two key components: first, we need a way to identify neurons that are not contributing to the learning process, and second we need a way to balance the network size and the generalization capability for the main task. We introduce a new Filter layer that takes care of *deactivating* neurons. We also modify existing loss functions to incorporate a new term that takes care of balancing network size and generalization capability appropriately.

Filter Layers: Filter layers have weights in the range $[0, +\infty]$ and are placed after linear and convolutional layers. The *Filter Layer* takes an input of size $(B \times C \times D_1 \times \dots \times D_n)$, where B is the batch size, C the number of features (or channels, in the case of convolutional layers), and D any additional dimension. This structure makes it compatible with fully connected layers with $n = 0$ or convolutional layers with $n = 2$. Their crucial property is a parameter $\theta \in \mathbb{R}^C$, defined as follows:

$$Filter(I; \theta) = I \circ \max(0, \theta) \quad (1)$$

Where \circ is the pointwise multiplication, and θ is expanded in all dimensions except the second one **Sam: why not the second one? what is the second one? C?** to match the input size. It is easy to see that if for any k , if $\theta_k \leq 0$, the k^{th} feature/channel will forever be 0. When this happens, we say the Filter layer deactivates the neuron. **Sam: Not easy to see this because you haven't told me that theta only decays over time?** These disabled neurons/channels can be removed from the network without changing its output. Before explaining how that is achieved, we explain next how the weights of the Filter Layer are adjusted during training. **Sam: How are filter weights initialized?**

Training Procedure: Once Filter layers are placed in a network, we could train the network directly using our standard loss function, and we could achieve performance equivalent to a normal neural network. However, our goal is to find the smallest network with reasonable performance. We achieve that by introducing sparsity in the parameters of the *Filter Layers*, thus forcing the deactivation of neurons **Sam: by decaying the value of theta** To obtain this sparsity, we simply redefine the loss function:

$$L'(x, y; \theta) = L(x, y) + \lambda |\theta| \quad (2)$$

The additional term $\lambda |\theta|$ introduces sparsity (see Lasso loss [?]). **Sam: Say lambda is in [0,1]?** The λ parameter adjusts how aggressively the network deactivates neurons, with larger values indicating more aggressive deactivation. **Sam: Give some intuitive statement about how this loss function causes theta to decay in irrelevant neurons, while preserving the value of important neurons.** Because the entries in θ with a value of 0 correspond to dead neurons, λ effectively controls the number of neurons/channels in the entire network. Next, we explain how to implement ShrinkNets efficiently.

Dynamic Network Resizing It is possible to reduce the overhead of the training process by removing neurons as soon as they become deactivated by θ going to 0. To do this, we implemented a *neural garbage collection* mechanism which prunes deactivated neurons on-the-fly, reducing the processing time and memory overhead. To support this feature, it is crucial to understand the information flow between neurons and layers in the neural network. We achieve this by representing such information flow as a graph, in which edges indicate what action each node, i.e., neuron, must take. **Sam: I can't really understand what this means – need to add a bit more detail.**

3 EVALUATION

We implemented ShrinkNets, including the *garbage collection* mechanism on top of *PyTorch* [?], and we have made the software open source.¹ We used our implementation to evaluate two crucial aspects of ShrinkNets, which we present next.

¹The code is available as Python/PyTorch library on <http://github.com/mitdbg/fastdeepnets>




Figure 1: Evolution of the number of hidden units and loss over time on the MNIST dataset for different λ values

Sam: Move legend inside top of figure, to save space.

3.1 Does ShrinkNet converge?

We trained a one hidden layer neural network with one filter layer to control the number of hidden units. We initialized the models with 10000 neurons and trained them on MNIST [?] using different values for λ . Figure Figure 1 shows that λ works as a proxy for the network size, with bigger λ implying smaller networks. More importantly, the figure helps us confirm that despite the regular spikes—which are caused by the dynamic disabling of neurons—ShrinkNets eventually converge.

ra: stopped here

3.2 Does ShrinkNet find good networks?

One of our main goals is to show that ShrinkNets quickly finds neural network architectures that perform well. To show this, we ran a simple experiment where we compare ShrinkNets to other network tuning methods, using random search over the hyper-parameters that control the network size (λ in the case of ShrinkNets). Despite the simple nature of this experiment, we believe that the results generalize to more complex search methods, many of which will actually benefit dimensionality of the hyper-parameter space in the ShrinkNets approach.

To isolate the effect of ShrinkNets we will consider two very simple architectures: multi-layer perceptrons (with three hidden layers) and the *LeNet-5* model [?]. For this experiment we assume we have no prior information about the size except an upper bound of 50 channels per convolutional layer and 5000 neurons for fully connected layers. We try to find the appropriate size of the network using classical network-size optimization methods and ShrinkNets. For the classical methods we use: i) purely random selection of network size parameters (“Fixed”), as well as hyper-parameter regularization over the network size parameters using L_2 penalty [?] (“Fixed+ L_2 ”). In the experiment, we randomly sample parameters for each class of network and train them for the same number of

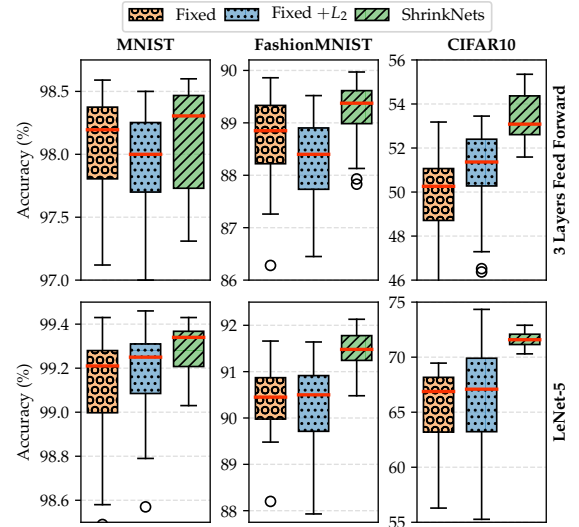


Figure 2: Distributions of the testing accuracy for different training methods, datasets and architectures using random search

epochs (100 for MLPs and 200 for CNNs). We picked the epoch that performed the best on the validation set and evaluated it on the testing set. We repeat the process for 50 iterations. This evaluation was done on MNIST [?], FashionMNIST [?] and CIFAR10 [?]. The distribution of the testing accuracy we obtained is summarized in Figure 2.

We can see that in all cases the distributions for ShrinkNets have substantially higher medians and are less variable than the other architectures (the best model is not always found by ShrinkNets, but this is likely an artifact of our randomized search process). This shows that ShrinkNets can find better networks with fewer training iterations than existing methods. **Sam:** Note something about how we expect this to bear out even with methods other than random search?

4 RELATED WORK

There are many methods in the literature that aim to simplify network structure. Most of them focuses on removing connections (eg: [?], [?]). However, most hardware (GPU’s, TPUS [?], Tensor-Cores) do not really leverage sparsity in the connections efficiently (as explained in [?]). On the other side, ShrinkNets and some others [?] and [?] try to remove entire neurons instead of connections. This is very useful because it reduces the size of the matrices, producing speed-up on every device. ShrinkNets improves on [?] because it removes neurons during training, speeding up the rest of the process. Further it outperforms [?] on convergence speed and because it does not need to change the optimizer. To the best of our knowledge this is also the first work that tries to learn the number of channels in a convolutional neural network.

5 CONCLUSION

In this paper we presented a novel technique to guess a reasonable network size based on single parameter λ that control the tradeoff between loss and network size. We demonstrated that ShrinkNets works both on fully connected networks and convolutional neural

networks. Although the initial results are promising, there are many additional avenues to explore. In the current implementation we only “learn” the number of features (neurons or channels). We plan to augment this with a dynamic number of layers as seen in [?] to be able to determine the entire architecture. Further, as shown in Figure 1 that the loss temporarily suffers from the removal of neurons. It is likely that the loss would be more stable if the number of neurons converged faster or neurons disappeared more slowly. For this reason we plan to explore proximal gradient methods to optimize the filter vectors and/or randomize neuron removals. Finally, during our evaluation we picked small datasets mainly to be able to train many models and have statistically significant distributions. We plan to verify that our approach see if it generalizes to bigger datasets and other architectures like ResNet [?], which is possible with small modifications to our existing code base.